

Wine Recognition Dataset Classification

Classification

- Covers classification based model
- Built various classification models based on different algorithm
- All models are based on Wine recognition dataset

Imports and Settings

Here are all the settings and imports used in the project.

```
17
18
19 import numpy as np
20 import pandas as pd
21 # Visualization Libraries
22 import matplotlib.pyplot as plt
23 import seaborn as sns
24 from sklearn import datasets # for using built-in datasets
25 from sklearn import metrics # for checking the model accuracy
26 #To plot the graph embedded in the notebook
27 import matplotlib.pyplot as plt
28 from pandas.plotting import parallel_coordinates
29 from sklearn.model_selection import train_test_split
30 # importing the necessary package to use the classification algorithm
31 from sklearn import svm #for Support Vector Machine (SVM) Algorithm
32 # importing the necessary package to use the classification algorithm
33 from sklearn.tree import DecisionTreeClassifier #for using Decision Tree Algoithm
34 # importing the necessary package to use the classification algorithm
35 from sklearn.neighbors import KNeighborsClassifier # for K nearest neighbours
36 # importing the necessary package to use the classification algorithm
37 from sklearn.linear_model import LogisticRegression # for Logistic Regression algorithm
38 # importing the necessary package to use the classification algorithm
39 from sklearn.naive_bayes import GaussianNB
40
```

Load Dataset: Wine recognition

- We will load the wine recognition dataset from the scikit-learn library.
- The objective is to predict the class of wine using the given features.

Input:

```
27
28 from sklearn.datasets import load_wine
29 wineData = load_wine()
30
```

Explore the Dataset

- sklearn returns Dictionary-like object.

Input:

```
31
32 print(wineData.keys())
33
```

Output:

```
dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR', 'feature_names'])
```

Input:

```
33
34 # Let's print the description of wine dataset
35 print(wineData.DESCR)
36
```

Output:

```
.. _wine_dataset:
```

Wine recognition dataset

****Data Set Characteristics:****

:Number of Instances: 178 (50 in each of three classes)

:Number of Attributes: 13 numeric, predictive attributes and the class

:Attribute Information:

- Alcohol
- Malic acid
- Ash
- Alcalinity of ash
- Magnesium
- Total phenols
- Flavanoids
- Nonflavanoid phenols
- Proanthocyanins
- Color intensity
- Hue
- OD280/OD315 of diluted wines

- Proline

- class:

- class_0

- class_1

- class_2

:Summary Statistics:

	Min	Max	Mean	SD
Alcohol:	11.0	14.8	13.0	0.8
Malic Acid:	0.74	5.80	2.34	1.12
Ash:	1.36	3.23	2.36	0.27
Alcalinity of Ash:	10.6	30.0	19.5	3.3
Magnesium:	70.0	162.0	99.7	14.3
Total Phenols:	0.98	3.88	2.29	0.63
Flavanoids:	0.34	5.08	2.03	1.00
Nonflavanoid Phenols:	0.13	0.66	0.36	0.12
Proanthocyanins:	0.41	3.58	1.59	0.57
Colour Intensity:	1.3	13.0	5.1	2.3
Hue:	0.48	1.71	0.96	0.23
OD280/OD315 of diluted wines:	1.27	4.00	2.61	0.71
Proline:	278	1680	746	315

:Missing Attribute Values: None

:Class Distribution: class_0 (59), class_1 (71), class_2 (48)

:Creator: R.A. Fisher

:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)

:Date: July, 1988

This is a copy of UCI ML Wine recognition datasets.

<https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data>

The data is the results of a chemical analysis of wines grown in the same region in Italy by three different cultivators. There are thirteen different measurements taken for different constituents found in the three types of wine.

Original Owners:

Forina, M. et al, PARVUS -

An Extendible Package for Data Exploration, Classification and Correlation.

Institute of Pharmaceutical and Food Analysis and Technologies,

Via Brigata Salerno, 16147 Genoa, Italy.

Citation:

Lichman, M. (2013). UCI Machine Learning Repository

[<https://archive.ics.uci.edu/ml>]. Irvine, CA: University of California,

School of Information and Computer Science.

.. topic:: References

(1) S. Aeberhard, D. Coomans and O. de Vel,

Comparison of Classifiers in High Dimensional Settings,

Tech. Rep. no. 92-02, (1992), Dept. of Computer Science and Dept. of

Mathematics and Statistics, James Cook University of North Queensland.

(Also submitted to Technometrics).

The data was used with many others for comparing various classifiers. The classes are separable, though only RDA has achieved 100% correct classification.

(RDA : 100%, QDA 99.4%, LDA 98.9%, 1NN 96.1% (z-transformed data))

(All results using the leave-one-out technique)

(2) S. Aeberhard, D. Coomans and O. de Vel,

"THE CLASSIFICATION PERFORMANCE OF RDA"

Tech. Rep. no. 92-01, (1992), Dept. of Computer Science and Dept. of Mathematics and Statistics, James Cook University of North Queensland.

(Also submitted to Journal of Chemometrics).

- As we can see from the dataset description that there is no missing values.
- And, it contains equal number of samples for each of the class of wine.

Input:

```
36
37 #Let's print the data (features matrix) of wine dataset
38 print(wineData.data)
39
```

Output:

```
[[1.423e+01 1.710e+00 2.430e+00 ... 1.040e+00 3.920e+00 1.065e+03]
 [1.320e+01 1.780e+00 2.140e+00 ... 1.050e+00 3.400e+00 1.050e+03]
 [1.316e+01 2.360e+00 2.670e+00 ... 1.030e+00 3.170e+00 1.185e+03]
 ...
 [1.327e+01 4.280e+00 2.260e+00 ... 5.900e-01 1.560e+00 8.350e+02]
 [1.317e+01 2.590e+00 2.370e+00 ... 6.000e-01 1.620e+00 8.400e+02]
 [1.413e+01 4.100e+00 2.740e+00 ... 6.100e-01 1.600e+00 5.600e+02]]
```

Input:

```
39
40 # Let's check the shape of features matrix
41 print(wineData.data.shape)
42
```

Output:

(178, 13)

Input:

```
42
43 # Let's print the feature names
44 print(wineData.feature_names)
45
```

Output:

['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins', 'color_intensity', 'hue', 'od280/od315_of_diluted_wines', 'proline']

Input:

```
45  
46 # Let's print the target vector  
47 print(wineData.target)  
48
```

Output:

[illegible]

Input:

```
48
49 # Let's check the shape of target
50 print(wineData.target.shape)
51
```

Output:

(178,)

Input:

```
51
52 #Let's print the target class/species names
53 print(wineData.target_names)
54
```

176	13.17	2.59	2.37	20.0	120.0	...	1.46	9.30	0.60
1.62	840.0								

0

..
... ...

173 13.71 5.65 2.45 20.5 95.0 1.68 ... 1.06 7.70 0.64
1.74 740.0

2

174 13.40 3.91 2.48 23.0 102.0 1.80 ... 1.41 7.30 0.70
1.56 750.0

2

175 13.27 4.28 2.26 20.0 120.0 1.59 ... 1.35 10.20 0.59
1.56 835.0

2

176 13.17 2.59 2.37 20.0 120.0 1.65 ... 1.46 9.30 0.60
1.62 840.0

2

177 14.13 4.10 2.74 24.5 96.0 2.05 ... 1.35 9.20 0.61
1.60 560.0

2

[178 rows x 14 columns]

Target species names:

- 0 = 'class_0'
- 1 = 'class_1'
- 2 = 'class_2'

Input:

```

61
62 # replace the target values with class names
63 wine_df['recog'] = wine_df['recog'].replace([0, 1, 2], ['class_0', 'class_1', 'class_2'])
64 print(wine_df)
65

```

Output:

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	od280/od315_of_diluted_wines	proline	recog
0	14.23	1.71	2.43		15.6	127.0	2.80	3.06				0.28		
	2.29	5.64	1.04		3.92	1065.0	class_0							
1	13.20	1.78	2.14		11.2	100.0	2.65	2.76				0.26		
	1.28	4.38	1.05		3.40	1050.0	class_0							
2	13.16	2.36	2.67		18.6	101.0	2.80	3.24				0.30		
	2.81	5.68	1.03		3.17	1185.0	class_0							
3	14.37	1.95	2.50		16.8	113.0	3.85	3.49				0.24		
	2.18	7.80	0.86		3.45	1480.0	class_0							
4	13.24	2.59	2.87		21.0	118.0	2.80	2.69				0.39		
	1.82	4.32	1.04		2.93	735.0	class_0							
...
...
173	13.71	5.65	2.45		20.5	95.0	1.68	0.61				0.52		
	1.06	7.70	0.64		1.74	740.0	class_2							
174	13.40	3.91	2.48		23.0	102.0	1.80	0.75				0.43		
	1.41	7.30	0.70		1.56	750.0	class_2							
175	13.27	4.28	2.26		20.0	120.0	1.59	0.69				0.43		
	1.35	10.20	0.59		1.56	835.0	class_2							
176	13.17	2.59	2.37		20.0	120.0	1.65	0.68				0.53		
	1.46	9.30	0.60		1.62	840.0	class_2							
177	14.13	4.10	2.74		24.5	96.0	2.05	0.76				0.56		
	1.35	9.20	0.61		1.60	560.0	class_2							

[178 rows x 14 columns]

Exploratory Data Analysis

Input:

```
70
71 # Return numerical summary of each attribute of wine
72 print(wine_df.describe())
73
```

Output:

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	...
	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	od280/od315_of_diluted_wines	proline	
count	178.000000	178.000000	178.000000	178.000000	178.000000	178.000000	...
	178.000000	178.000000	178.000000	178.000000	178.000000	178.000000	
mean	13.000618	2.336348	2.366517	19.494944	99.741573	2.295112	...
	0.361854	1.590899	5.058090	0.957449	2.611685	746.893258	
std	0.811827	1.117146	0.274344	3.339564	14.282484	0.625851	...
	0.124453	0.572359	2.318286	0.228572	0.709990	314.907474	
min	11.030000	0.740000	1.360000	10.600000	70.000000	0.980000	...
	0.130000	0.410000	1.280000	0.480000	1.270000	278.000000	
25%	12.362500	1.602500	2.210000	17.200000	88.000000	1.742500	...
	0.270000	1.250000	3.220000	0.782500	1.937500	500.500000	
50%	13.050000	1.865000	2.360000	19.500000	98.000000	2.355000	...
	0.340000	1.555000	4.690000	0.965000	2.780000	673.500000	
75%	13.677500	3.082500	2.557500	21.500000	107.000000	2.800000	...
	0.437500	1.950000	6.200000	1.120000	3.170000	985.000000	
max	14.830000	5.800000	3.230000	30.000000	162.000000	3.880000	...
	0.660000	3.580000	13.000000	1.710000	4.000000	1680.000000	

[8 rows x 13 columns]

Input:

```
66
67 # let's check number of samples for each class of wine
68 print(wine_df.groupby('recog').size())
69
```

Output:

recog

class_0 59

class_1 71


class_2 48

dtype: int64

Input:

```
69
70 # let's visualise the number of samples for each class with count plot
71 sns.countplot(x='recog', data=wine_df)
72 plt.show()
73
```

Output:




- Next, let's make a correlation matrix to quantitatively examine the relationship between variables.
- If there are features and many of the features are highly correlated, then training an algorithm with all the features will reduce the accuracy. Thus features selection should be done carefully. This dataset has less features but still we will see the correlation.
- The correlation matrix can be formed by using the corr function from the pandas library.
- The correlation coefficient ranges from -1 to 1 . If the value is close to 1 , it means that there is a strong positive correlation between the two variables. When it is close to -1 , the variables have a strong negative correlation.
- Then, we will use the heatmap function from the seaborn library to plot the correlation matrix.

Input:

```
79
80 # changing the figure size
81 plt.figure(figsize = (9, 6))
82 # "annot = True" to print the values inside the square
83 sns.heatmap(data=correlation_matrix, annot=True)
84 plt.show()
85
```


Output:



Input:

```
85
86 # Steps to remove redundant values
87 # Return a array filled with zeros
88 mask = np.zeros_like(correlation_matrix)
89 # Return the indices for the upper-triangle of array
90 mask[np.triu_indices_from(mask)] = True
91 # changing the figure size
92 plt.figure(figsize = (9, 6))
93 # "annot = True" to print the values inside the square
94 sns.heatmap(data=correlation_matrix, annot=True, mask=mask)
95 plt.show()
96
```

Output:




Input:

```

96
97 # let's create pairplot to visualise the data for each pair of attributes
98 sns.pairplot(wine_df, hue="recog", height = 2, palette = 'colorblind')
99 plt.show()
100

```

Output:




- As we can see from the above figure, wine classes are pretty close to each other.
- For this dataset, another useful visualization plot is parallel coordinate, which represents each row as a line.
- A parallel plot allows to compare the feature of several individual observations (series) on a set of numeric variables. Interestingly, Pandas is probably the best way to plot a parallel coordinate plot with python.

Input:

```
101
102 parallel_coordinates(wine_df, "recog", color = ['blue', 'red', 'green'])
103 plt.show()
104
```

Output:



Create Features Matrix & Target Variable

Input:

```
105
106 x = wine_df[['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins',
107             'color_intensity', 'hue', 'od280/od315_of_diluted_wines', 'proline']]
108 print(x)
```

Output:

alcohol malic_acid ash alcalinity_of_ash magnesium total_phenols flavanoids
nonflavanoid_phenols proanthocyanins color_intensity hue od280/od315_of_diluted_wines
proline

0	14.23	1.71	2.43		15.6	127.0	2.80	3.06	0.28
	2.29	5.64	1.04		3.92	1065.0			
1	13.20	1.78	2.14		11.2	100.0	2.65	2.76	0.26
	1.28	4.38	1.05		3.40	1050.0			
2	13.16	2.36	2.67		18.6	101.0	2.80	3.24	0.30
	2.81	5.68	1.03		3.17	1185.0			
3	14.37	1.95	2.50		16.8	113.0	3.85	3.49	0.24
	2.18	7.80	0.86		3.45	1480.0			
4	13.24	2.59	2.87		21.0	118.0	2.80	2.69	0.39
	1.82	4.32	1.04		2.93	735.0			
...
...	...								
173	13.71	5.65	2.45		20.5	95.0	1.68	0.61	0.52
	1.06	7.70	0.64		1.74	740.0			
174	13.40	3.91	2.48		23.0	102.0	1.80	0.75	0.43
	1.41	7.30	0.70		1.56	750.0			
175	13.27	4.28	2.26		20.0	120.0	1.59	0.69	0.43
	1.35	10.20	0.59		1.56	835.0			
176	13.17	2.59	2.37		20.0	120.0	1.65	0.68	0.53
	1.46	9.30	0.60		1.62	840.0			

177	14.13	4.10	2.74	24.5	96.0	2.05	0.76	0.56
1.35	9.20	0.61	1.60	560.0				

[178 rows x 13 columns]

Input:

```
108
109 y = wine_df['recog']
110 print(y)
111
```

Output:

```
0    class_0
1    class_0
2    class_0
3    class_0
4    class_0
...
173   class_2
174   class_2
175   class_2
176   class_2
177   class_2
```

Name: recog, Length: 178, dtype: object

Split the dataset

Input:

```
111
112 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 16)
113 print("X_train shape: ", X_train.shape)
114 print("X_test shape: ", X_test.shape)
115 print("y_train shape: ", y_train.shape)
116 print("y_test shape: ", y_test.shape)
117
```

Output:

X_train shape: (124, 13)

X_test shape: (54, 13)

y_train shape: (124,)

y_test shape: (54,)

Create Model: Support Vector Machine (SVM)

Input:

```
118 model_svm = svm.SVC() #select the algorithm
119 model_svm.fit(X_train, y_train) #train the model with the training dataset
120 y_prediction_svm = model_svm.predict(X_test) # pass the testing data to the trained model
121 # checking the accuracy of the algorithm.
122 # by comparing predicted output by the model and the actual output
123 score_svm = metrics.accuracy_score(y_prediction_svm, y_test).round(4)
124 print("-----")
125 print('The accuracy of the SVM is: {}'.format(score_svm))
126 print("-----")
127 # save the accuracy score
128 score = set()
129 score.add(('SVM', score_svm))
130
131
```

Output:

The accuracy of the SVM is: 0.6111

Create Model: Decision Tree

Input:

```
132 model_dt = DecisionTreeClassifier(random_state=4)
133 model_dt.fit(X_train, y_train) #train the model with the training dataset
134 y_prediction_dt = model_dt.predict(X_test) #pass the testing data to the trained model
135 # checking the accuracy of the algorithm.
136 # by comparing predicted output by the model and the actual output
137 score_dt = metrics.accuracy_score(y_prediction_dt, y_test).round(4)
138 print("-----")
139 print('The accuracy of the DT is: {}'.format(score_dt))
140 print("-----")
141 # save the accuracy score
142 score = set()
143 score.add(('DT', score_dt))
144
145
```

Output:

The accuracy of the DT is: 0.9444

Create Model: K Nearest Neighbours (KNN)

Input:

```
146
147 from sklearn.linear_model import LogisticRegression # for Logistic Regression algorithm
148 model_knn = KNeighborsClassifier(n_neighbors=3) # 3 neighbours for putting the new data into a class
149 model_knn.fit(X_train, y_train) #train the model with the training dataset
150 y_prediction_knn = model_knn.predict(X_test) #pass the testing data to the trained model
151 # checking the accuracy of the algorithm.
152 # by comparing predicted output by the model and the actual output
153 score_knn = metrics.accuracy_score(y_prediction_knn, y_test).round(4)
154 print("-----")
155 print('The accuracy of the KNN is: {}'.format(score_knn))
156 print("-----")
157 # save the accuracy score
158 score = set()
159 score.add(('KNN', score_knn))
160
```

Output:

The accuracy of the KNN is: 0.7037

Create Model: Logistic Regression

Input:

```
162
163 model_lr = LogisticRegression()
164 model_lr.fit(X_train, y_train) #train the model with the training dataset
165 y_prediction_lr = model_lr.predict(X_test) #pass the testing data to the trained model
166 # checking the accuracy of the algorithm.
167 # by comparing predicted output by the model and the actual output
168 score_lr = metrics.accuracy_score(y_prediction_lr, y_test).round(4)
169 print("-----")
170 print('The accuracy of the LR is: {}'.format(score_lr))
171 print("-----")
172 # save the accuracy score
173 score = set()
174 score.add(('LR', score_lr))
175
```

Output:

The accuracy of the LR is: 0.9444

Input:

```

177 model_nb = GaussianNB()
178 model_nb.fit(X_train, y_train) #train the model with the training dataset
179 y_prediction_nb = model_nb.predict(X_test) #pass the testing data to the trained model
180 # checking the accuracy of the algorithm.
181 # by comparing predicted output by the model and the actual output
182 score_nb = metrics.accuracy_score(y_prediction_nb, y_test).round(4)
183 print("-----")
184 print('The accuracy of the NB is: {}'.format(score_nb))
185 print("-----")
186 # save the accuracy score
187 score = set()
188 score.add(('NB', score_nb))
189
190

```

Output:

The accuracy of the NB is: 1.0

- Best accuracy

Compare Accuracy Score of Different Models

The accuracy scores of different Models:

('NB', 1.0)

('LR', 0.9444)

('SVM', 0.6111)

('KNN', 0.7037)

('DT', 0.9444)

- Here we can see that all the models accuracy score are not good.
- Only accuracy score above 0.9 is good.
- NB model gave the prefect accuracy